

ECE 239AS Problem Set 3 - (VIGNESH NAGARAJAN)

Imports

```
import os
from google.colab import drive

import numpy as np
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from PIL import Image

import torch
import torch.nn as nn
import torch.nn.functional as F

from torch.utils.data import DataLoader, random_split
from torchvision import transforms

drive.mount('/content/drive')
%cd '/content/drive/MyDrive/239AS/Assignment-3'

Mounted at /content/drive
/content/drive/MyDrive/239AS/Assignment-3

# !gdown https://drive.google.com/uc?id=1uPSqGZdmZt-
# b5YLMWscad6ST3MLBNESk
# !unzip sol.zip

def check_code_correctness(test_function, *test_in, test_out):
    """Checks if the given function behaves as expected

    Args:
        test_function: Function to test
        *test_in: The sample inputs to test on (can be multiple)
        test_out: The expected output of the function given the
        specified input

    Returns:
        True if the function behaves as expected, False otherwise (wrong
        answer or error)
    """
    try:
        # Using *test_in to unpack the arguments and pass to the
        test_function
```

```

        student_outputs = test_function(*test_in)
        # Ensure student_outputs is a tuple for easier comparison
        if not isinstance(student_outputs, tuple):
            student_outputs = (student_outputs,)
        except NotImplementedError as err:
            print("Please implement and remove 'raise
NotImplementedError'")
            return False
        except RuntimeError as err:
            print("Please make sure you have the right dimensions and
type")
            return False
        except:
            print("An exception occurred: could not compute output")
            return False

    try:
        for student_out, test_out in zip(student_outputs, test_outs):
            if not np.allclose(student_out, test_out):
                print("Test failed, student output does not match test
output")
                return False
        except TypeError as err:
            print("Please make sure your function outputs the correct
type")
            return False
        except:
            print("An exception occurred: could not check output")
            return False

    return True

```

0 Introduction to NeRF and Pytorch

0.1 NeRF Introduction

0.1.1 Intuition of NeRF



Before we implement NeRF, let's introduce the idea and intuition behind it.

Imagine you have an object (a phone, an article of clothing, a yellow lego tractor) in the real world that you want display online. One's first intuition may be to take a picture of the front of the item and then post that. Others may take photos of the item at different angles. Some may take a video around the object to allow even more angles to view it. However, each of these techniques fall short in the same way: given a finite amount of pictures/frames, you can only get a finite amount of ways to view it.

This is where Neural Radiance Fields (NeRF) comes in. The main application of NeRF is novel view synthesis -- being able to create new poses or views of an object. This has major implications: given a finite amount of pictures/frames, you can get an infinite amount of ways to view it.

While novel view synthesis may not seem that important on its own, it has many implications in computer vision and computational imaging. When humans look at objects, we have an understanding of what different views of the object are (back, top, side) based on our previous experiences with it. Thus, if we want to create a machine that can interpret and understand 3D real world scenes, we must incorporate a way for the machine to understand not observed views of objects.

Another implication of this is surface reconstruction, which was introduced in the paper "NeuS: Learning Neural Implicit Surfaces by Volume Rendering for Multi-view Reconstruction". Given 2D images of a scene, this neural network is able to create an SDF that we can extract a mesh from using marching cubes. This mesh can then be imported to Blender.

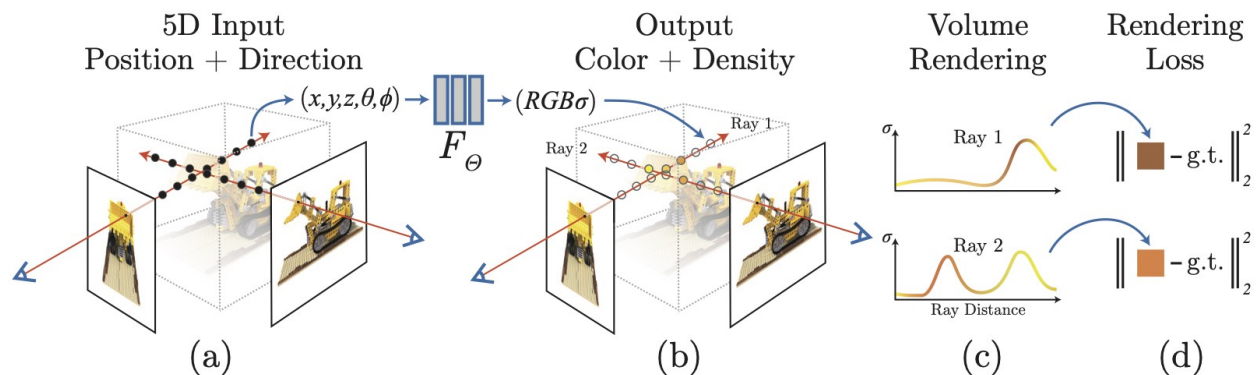
High-Level Overview of NeRF

So how does NeRF work? We will use the following analogy throughout this notebook:

Imagine you are trying to paint a picture, but instead of using a brush, you shine light through a foggy room and capture the colors that come out on the other side. The more fog (or "stuff") there is in a particular spot, the more it affects the color of the light. If there's a lot of fog, we might not even see the light that comes from behind it. To paint a specific view of the scene, you anchor your flashlight somewhere and shine it into the fog. The light that comes out of the fog is dependent on the color and fog density at different points along the flashlight's beams. You can create different views of the fog by shining the flashlight at different angles and positions (you can't change the fog).

Now, let's apply this analogy to the yellow lego tractor. Here, our scene is the lego tractor. Instead of using a flashlight, we use a camera to "shoot" rays that travel across the scene. As we only have pictures, we use a neural network to "simulate" what the rgb and density/opacity of the scene is at different points. We then use alpha blending to turn these rgb and density/opacity values into an image. We can train the neural network by minimizing the difference between an image of the scene at a specific viewing position and angle and the reconstructed image from the neural network given the viewing parameters.

The end result is the following function: given a camera angle and position, we can produce an image of the scene from that angle.



0.2 Pytorch Introduction

0.2.1 Numpy and Pytorch

Many functions from pytorch are identical to numpy's versions (except that they output tensors instead of arrays). Try it out on linspace. (Hint: Use torch.linspace)

```
def student_linspace(start, end, num):
    return torch.linspace(0,1,4)

print(np.linspace(0,1,4))
print(student_linspace(0,1,4))

[0.          0.33333333 0.66666667 1.          ]
tensor([0.0000, 0.3333, 0.6667, 1.0000])
```

One function you should know is torch.from_numpy(). This will convert a numpy array into a torch tensor.

```
def student_tensor(t):
    return torch.from_numpy(t)


print(np.array([1,2,3]))
print(student_tensor(np.array([1,2,3])))

[1 2 3]
tensor([1, 2, 3])
```

0.2.2 Indexing

A key component of both numpy and pytorch are high-dimensional tensors and specifically indexing into them. As an example, lets look at a HxWx3 tensor representing an image. We'll start by loading in an image.

```
image =  
transforms.functional.pil_to_tensor(Image.open("sol/cornell_box.png"))  
.permute((1,2,0))  
plt.imshow(image)  
image.shape  
  
torch.Size([300, 300, 3])
```




The 3 in the shape refers to the 3 different RGB channels of the image. What if we want to get only the R channel? To do this, we can use indexing syntax like shown below. We have 3 things inside the brackets because our original tensor has 3 dimensions. If we put ":" for a dimension, that means we want everything in that dimension, and if we put a number of range of numbers, that means we want that range.

To get R, we want everything in the height and width dimensions but only the first value in the color dimension:

```
red = image[:, :, 0]  
plt.imshow(red)
```

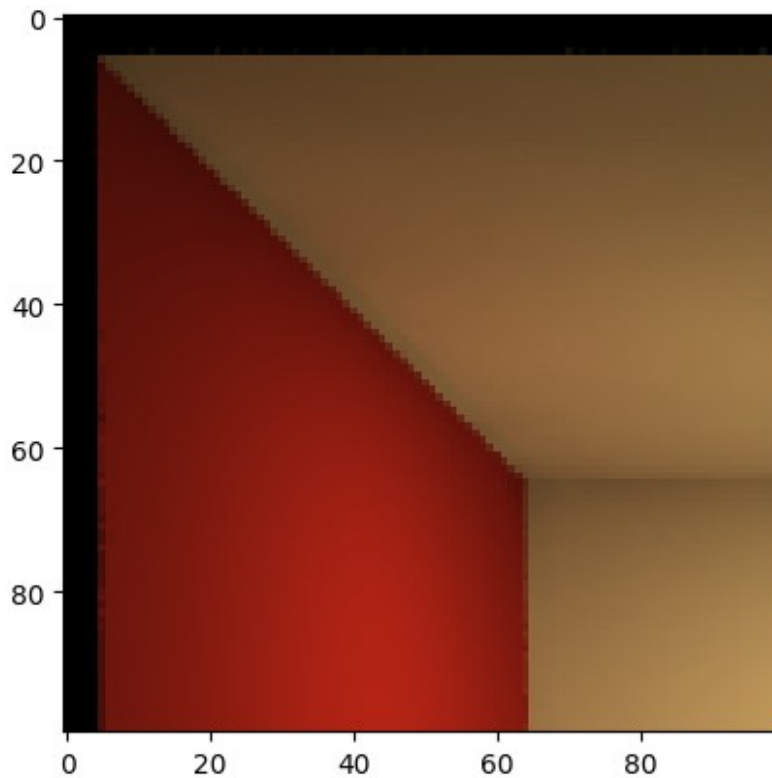
```
<matplotlib.image.AxesImage at 0x795a01988070>
```



What if instead we wanted the top left 100x100 corner in all 3 colors? How would you do that?

```
top_left = image[:,100,:100,:]
plt.imshow(top_left)
```

```
<matplotlib.image.AxesImage at 0x7959fe7d5ab0>
```



If you want to add a dimension to a tensor, you can use ".unsqueeze(dim)" and the corresponding ".squeeze(dim)" to undo it. An example:

```
test_tensor = torch.zeros((4,4))
print(f"Shape at start: {test_tensor.shape}")

test_tensor = test_tensor.unsqueeze(2)
print(f"Shape after unsqueeze: {test_tensor.shape}")

test_tensor = test_tensor.squeeze(2) # What happens if I try to
squeeze out dim 0 or 1 instead?
print(f"Shape after squeeze: {test_tensor.shape}")

Shape at start: torch.Size([4, 4])
Shape after unsqueeze: torch.Size([4, 4, 1])
Shape after squeeze: torch.Size([4, 4])
```

You can also index using another array, either an integer or boolean array:

```
x = torch.arange(10) - 5
select_idx = torch.tensor([0,2,5,6,8])
print(f"X: {x}")
print(f"selected: {x[select_idx]}")
```

```

select_idx = (x > 0)
print(f"selected positive only: {x[select_idx]}")

X: tensor([-5, -4, -3, -2, -1,  0,  1,  2,  3,  4])
selected: tensor([-5, -3,  0,  1,  3])
selected positive only: tensor([1, 2, 3, 4])

```

Another key function in torch/numpy is to combine multiple arrays in different ways (both mathematical and not).

To do mathematical operations, our arrays must have the same shape or have a shape that can be adapted to match. As an example:

```

x = torch.zeros((2,2,2))
y = torch.ones((2,2,2))

print(f"X: {x}")
print(f"Y: {y}")
print(f"X+Y: {x+y}")

z = torch.ones((2,2,1))
print(f"X+Z: {x+z}") # Torch will convert the "1" in dimension 2 to
                    # "2" automatically

t = torch.ones((2,2))
print(f"X+T: {x+t}") # Torch will add the extra dimension for you and
                    # duplicate

r = torch.ones((3,3))
try:
    print(f"X+R: {x+r}") # Wont work
except:
    print("X+R: Didn't work")

X: tensor([[[0., 0.],
            [0., 0.]],
          [[0., 0.],
            [0., 0.]])
Y: tensor([[[1., 1.],
            [1., 1.]],
          [[1., 1.],
            [1., 1.]])
X+Y: tensor([[[1., 1.],
            [1., 1.]],
          [[1., 1.],
            [1., 1.]])
X+Z: tensor([[[1., 1.],
            [1., 1.]])

```



```

        [1., 1.]],
        [[1., 1.],
         [1., 1.]])
X+T: tensor([[[1., 1.],
              [1., 1.]],
             [[1., 1.],
              [1., 1.]])
X+R: Didn't work

```

We may also want to combine matrices into larger matrices:

```

x = torch.zeros((2,2))
y = torch.ones((2,2))

z = torch.stack([x,y])
t = torch.cat([x,y])
t2 = torch.cat([x,y], dim=1)

print(f"X: {x}, shape: {x.shape}")
print(f"Y: {y}, shape: {y.shape}")
print(f"Z: {z}, shape: {z.shape}")
print(f"T: {t}, shape: {t.shape}")
print(f"T2: {t2}, shape: {t2.shape}")

X: tensor([[0., 0.],
           [0., 0.]]) shape: torch.Size([2, 2])
Y: tensor([[1., 1.],
           [1., 1.]]) shape: torch.Size([2, 2])
Z: tensor([[[0., 0.],
            [0., 0.]],
          [[1., 1.],
            [1., 1.]]) shape: torch.Size([2, 2, 2])
T: tensor([[0., 0.],
           [0., 0.],
           [1., 1.],
           [1., 1.]]) shape: torch.Size([4, 2])
T2: tensor([[0., 0., 1., 1.],
            [0., 0., 1., 1.]]) shape: torch.Size([2, 4])

```

0.2.3 Meshgrid

`torch.meshgrid()` is a useful function you will need when implementing NeRF. Intuitively, this function takes two ranges of numbers and then creates a 2D grid of all combinations of both ranges.

To test this, let's visualize two ranges:

If you want the x and y values of a HxW, use this function.


Note: There are two ways to index: "xy" and "ij". Try changing the indexing of the example and see how that affects your outputs. To learn more about meshgrid, you can read the following article: <https://www.geeksforgeeks.org/numpy-meshgrid-function/>

```
x = torch.tensor([1,2,3,4])
y = torch.tensor([1,2,3,4])

xx, yy = torch.meshgrid(x, y, indexing='xy')

fig, ax = plt.subplot_mosaic("""
                                AAAA
                                BBBB
                                CCCC
                                CCCC
                                CCCC
                                CCCC
                                DDDD
                                DDDD
                                DDDD
                                DDDD
                                """)

ax['A'].imshow(x.unsqueeze(0))
ax['B'].imshow(y.unsqueeze(0))
ax['C'].imshow(xx)
ax['D'].imshow(yy)
plt.show()
```



```
def example_meshgrid(H,W, indexing):
    return torch.meshgrid(torch.arange(H), torch.arange(W),
indexing=indexing)

fig, ax = plt.subplots(1,2)

x, y = example_meshgrid(3,4, indexing='ij')
print("*****ij indexing*****")
print(x)
x, y = x.flatten(), y.flatten()
# Create a colormap
colors = cm.rainbow(np.linspace(0, 1, len(x)))
ax[0].set_title('ij indexing')
for i in range(1, len(x)):
    ax[0].plot(x[i-1:i+1], y[i-1:i+1], color=colors[i])

x, y = example_meshgrid(3,4, indexing='xy')
print("*****xy indexing*****")
print(x)
print(y)

x, y = x.flatten(), y.flatten()

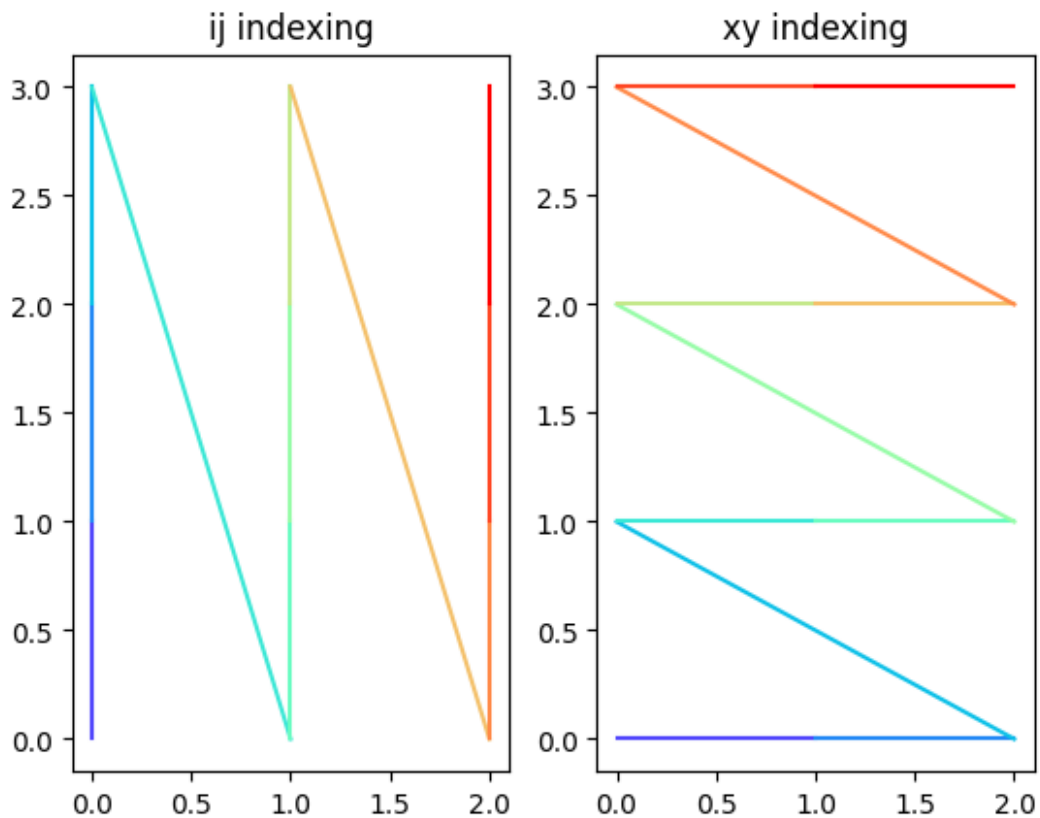
ax[1].set_title('xy indexing')
for i in range(1, len(x)):
```

```

ax[1].plot(x[i-1:i+1], y[i-1:i+1], color=colors[i])
plt.show()

*****ij indexing*****
tensor([[0, 0, 0, 0],
        [1, 1, 1, 1],
        [2, 2, 2, 2]])
*****xy indexing*****
tensor([[0, 1, 2],
        [0, 1, 2],
        [0, 1, 2],
        [0, 1, 2]])
tensor([[0, 0, 0],
        [1, 1, 1],
        [2, 2, 2],
        [3, 3, 3]])

```



As shown above, 'ij' indexing goes column by column, but 'xy' indexing goes row by row.

0.2.3 Cumulative Product

`torch.cumprod()` is another useful function you will need when implementing NeRF. Consider you have a tensor of values. This function will create a tensor whose *i*th value is the product of all elements before and including the *i*th value.

However, you will need a modified version of this function that allows the "exclusive" parameter. If this parameter is True, the output is a tensor whose ith value is the product of all elements before but NOT including the ith value.

```
def student_cumprod(t, exclusive=False):
    if not exclusive:
        return torch.cumprod(t,dim=-1)
    else:
        return torch.cumprod(t,dim=-1)/t

print("Input:", torch.tensor([1.5,2,3,4]))
print("Cumprod, inclusive:",
      student_cumprod(torch.tensor([1.5,2,3,4])))
print("Cumprod, exclusive:", student_cumprod(torch.tensor([1.5,2,3,4]),
                                              exclusive=True))

Input: tensor([1.5000, 2.0000, 3.0000, 4.0000])
Cumprod, inclusive: tensor([ 1.5000,  3.0000,  9.0000, 36.0000])
Cumprod, exclusive: tensor([1.0000, 1.5000, 3.0000, 9.0000])

a= torch.tensor([[1,2,3], [4,5,6]])
torch.cumprod(a,dim=1)/a

tensor([[ 1.,  1.,  2.],
        [ 1.,  4., 20.]])

torch.cumprod(a,dim=1)

tensor([[ 1,  2,  6],
        [ 4, 20, 120]])
```

1 Computing and Rendering Rays

In this section, please use pytorch instead of numpy. As we will be training a neural network, it will be easier to integrate your code instead of swapping between the two domains.

1.1 From Pixels to World Rays


1.1.1 Creating ray directions from pixels

Imagine rays being ejected from the camera to pierce each pixel in an image. The rays would start from the middle point of the sensor and hit the lens (where the image is). Each ray is being ejected from the camera, so they all point away in equal amounts in the z-direction. This means that the rays passing through the middle of the image are almost all (or exactly all) in the z-direction. Compute the directions of the rays for every pixel in the image.

Take a look at the figure below and try to derive the 3D position of a pixel in camera space.

Hint: The z component of all the rays are -1.

Hint: Imagine the line pointing from the camera center to one of the sides. What is the direction and length of the x and y components of the ray? Now consider a slightly smaller pyramid (one pixel removed from around the image). What is the direction and length of the x and y components of this new rays? How can we generalize this to get the x and y components of any ray in the image?



```
a= np.array([[1,2,3],[4,5,6]])
b = [[7,8,9], [10,11,12]]
a[... , np.newaxis, :]
```

```
array([[[[1, 2, 3]],
        [[4, 5, 6]]])
```

```
def student_compute_ray_directions(i,j,H,W,focal):
    """
    Compute ray directions given an image size and field of view.

    Parameters:
    - i (torch.Tensor, [H,W]): x index of pixels
    - j (torch.Tensor, [H,W]): y index of pixels
    - H (int): Height of the image (or viewport).
    - W (int): Width of the image (or viewport).
    - focal (float): The focal length of the camera.

    Returns:
    - dirs (torch.Tensor): Directions (x,y,z) of the rays for every
    pixel in the image. [H,W,3]
    """
    dirs = torch.stack([(i-W*0.5)/focal, -(j-H*0.5)/focal, -
    torch.ones_like(i)], -1 )
    return dirs
```

1.1.2 Compute World Rays

The directions and positions of rays we imagined in the last part are with respect to the camera. To get data that is generalizable and compatible to other camera positions and angles, we need to convert the rays from the camera's space to the world's space.

The c2w matrix is composed of a rotation matrix and a translation vector. The rotation matrix describes how to convert a ray from the camera's coordinate system to a generalizable coordinate system. The translation vector describes the camera's position in a generalizable space. Thus, the c2w matrix is unique to each pose an image is taken in. The c2w matrix is of shape 4x4. The c2w matrix is stored in the following way:

$$\begin{bmatrix} R & t \end{bmatrix}$$

Using the above information, compute rays_o -- a matrix containing the origins of each ray in the world's space -- and rays_d -- a matrix containing the direction of each ray in the world's space.

```
def student_compute_world_rays(dirs,c2w):
    """
    Convert the rays from camera space to world space.

    Parameters:
    - dirs (torch.Tensor, [H,W]): Directions of the rays for every
    pixel in the image.
    - c2w (torch.Tensor, [3,4]): A 3x4 camera-to-world transformation
    matrix.

    Returns:
    - rays_o (torch.Tensor, [H,W,3]): Origins of the rays in world
    space for every pixel in the image
    - rays_d (torch.Tensor, [H,W,3]): Directions of the rays in world
    space for every pixel in the image
    """
    rays_d = torch.sum(dirs[..., None, :] * c2w[:3, :3], -1)
    rays_o = torch.broadcast_to(c2w[:3, -1], np.shape(rays_d))

    return rays_o, rays_d
```

1.1.3 get_rays Function

Using the two previous two sections, create a function to compute a world origins and directions of rays from H, W, focal, and c2w.

```
def student_get_rays(H,W,focal,c2w):
    """
    Generate rays given an image size, focal length, and a camera-to-
    world transformation matrix.

    Parameters:
    - H (int): Height of the image (or viewport).
    - W (int): Width of the image (or viewport).
    - focal (float): The focal length of the camera.
    - c2w (torch.Tensor, [3,4]): A 3x4 camera-to-world transformation
    matrix.

    Returns:
```

```

- rays_o (torch.Tensor, [H,W,3]): Origins of the rays in world
space for every pixel in the image [H,W,3]
- rays_d (torch.Tensor, [H,W,3]): Directions of the rays in world
space for every pixel in the image [H,W,3]
"""
# Step 1: compute grid for x,y components
i, j = torch.meshgrid(torch.arange(W, dtype=torch.float32),
                      torch.arange(H, dtype=torch.float32),
indexing='xy')

# Step 2: For each pixel, compute ray directions
# Note: rays always point to camera (in z direction)
dirs = student_compute_ray_directions(i,j,H,W,focal)
# Step 3: Compute rays_o and rays_d
rays_o, rays_d = student_compute_world_rays(dirs,c2w)
return rays_o, rays_d

H_test = 32
W_test = 32
focal_test = 2
c2w_test = torch.eye(4)
rays_o_sol = torch.load('sol/rays_o.pt')
rays_d_sol = torch.load('sol/rays_d.pt')

check_code_correctness(student_get_rays,H_test,W_test,focal_test,c2w_test,
test_outs=[rays_o_sol,rays_d_sol])

True

```

1.2 Rendering Rays

Now that we have rays, we need to know what the colors along the rays is going to be. That will help us paint a view of the scene.

In particular, we will use a neural network that predicts the color and density of specific points. The process of us "painting" the picture will then look like this:

1. Sample points that we want to paint/render in our scene.
2. Use the neural network to determine the color and density at those points.
3. Find the weights associated with the density of each position.
4. Blend the colors using the weights to paint/render our scene.

1.2.1 Querying Points

Using the world rays we calculated in the last section, calculate where along the rays we should check the scene's color and density.

For each ray, we want to sample a set of linearly distanced points between the factors of 'near' and 'far'. These points are where we'll ask our neural network about the color and density of the scene.

Hint: Remember that we are sampling along multiples of the rays. Thus, each point should be the sum its rays origin as well as the direction multiplied by some factor.

```
def student_compute_query_points(rays_o, rays_d, near, far, N_samples,
rand=False):
    """
    Compute 3D query points along each ray.

    Inputs:
    - rays_o (torch.Tensor, [H,W,3]): Starting points of each ray.
    Think of this as where the camera is.
    - rays_d (torch.Tensor, [H,W,3]): Directions each ray is pointing
    towards.
    - near (float): How close to start looking along the ray.
    - far (float): How far to stop looking.
    - N_samples (int): How many points to check along each ray.
    - rand (bool): If true, jiggle the points a bit for smoother
    results.

    Returns:
    - pts (torch.Tensor, [H,W,N_samples,3]): 3D query points.
    - z_vals (torch.Tensor, [N_samples]): The distance along an
    arbitrary ray to each query point
    """

    # t_vals = torch.linspace(0.0, 1.0, N_samples).to(rays_o.device)
    z_vals = torch.linspace(near, far, N_samples).to(rays_o.device)
    #near * (1.0 - t_vals) + far * t_vals

    # perturb
    if rand:
        z_mids = 0.5 * (z_vals[1:] + z_vals[:-1])
        z_uppers = torch.cat([z_mids, z_vals[-1:]], dim=-1)
        z_lowers = torch.cat([z_vals[:1], z_mids], dim=-1)
        t_rand = torch.rand([N_samples]).to(rays_o.device)
        z_vals = z_lowers * (1.0 - t_rand) + z_uppers * t_rand
        z_vals = z_vals.expand(list(rays_o.shape[:-1]) + [N_samples])

    pts = rays_o[..., None, :] + z_vals[..., :, None] * rays_d[...,
None, :]

    return pts, z_vals

# depth_values = torch.linspace(near, far, N_samples).to(rays_o)
# if rand is True:
#     # ray_origins: (width, height, 3)
#     # noise_shape = (width, height, num_samples)
```

```

# noise_shape = list(rays_o.shape[:-1]) + [N_samples]
# # depth_values: (num_samples)
# depth_values = depth_values \
#     + torch.rand(noise_shape).to(rays_o) * (far
#     - near) / N_samples
# # (width, height, num_samples, 3) = (width, height, 1, 3) +
# (width, height, 1, 3) * (num_samples, 1)
# # query_points: (width, height, num_samples, 3)
# query_points = rays_o[..., None, :] + rays_d[..., None, :] *
depth_values[..., :, None]
# # TODO: Double-check that `depth_values` returned is of shape
# `(num_samples)`.
# return query_points, depth_values

```

1.2.2 Feed Query Points to Network

We will feed the 3D points we queried from the last section into our neural network. Remember, the neural network is like our painter. At each point we give it, the network will return the color and density at that point. However, our neural network cannot understand the inputs directly from the 3D coordinates. We will have to modify our inputs using a "positional encoding".

In this function, feed the points into the network and return the colors and opacities.

Hint: It can be very memory intensive to process every point at once! Thus, we have provided a "batchify" function to split up your inputs as you feed it into the network_fn. Hint: Reshape the output to be same shape as the input but with 4 channels (rgb, and density) instead of 3 at the end.

```

def student_query_network(network_fn, embed_fn, pts):
    """
    Ask the neural network about color and density for each point.

    Inputs:
    - network_fn (function): Our neural network.
    - embed_fn (function): The positional encoding function.
    - pts (torch.Tensor, [N_samples,3]): The points in space where we
    want to know about the scene's color and density.

    Returns:
    - raw (torch.Tensor, [H,W,N_samples,4]): The raw output of the
    neural network that returns the point's rgb and density.
    """
    pts = pts.to(device)
    def batchify(fn, chunk=1024*32):
        """Helper function to run the network in smaller batches."""
        return lambda inputs: torch.cat([fn(inputs[i:i+chunk]) for i
in range(0, inputs.shape[0], chunk)], 0)

    **STUDENT CODE**

```

```

height, width, n_samples, _ = pts.shape
pts_flattened = pts.reshape((-1,3))

batchified_embed_fn = batchify(embed_fn)
encoded_pts = batchified_embed_fn(pts_flattened)

batchified_fn = batchify(network_fn)
raw = batchified_fn(encoded_pts)
raw = raw.reshape((height, width, n_samples, -1))

return raw.to("cpu")

```

1.2.3 From Density to Weights

If a spot is very foggy/opaque, it'll affect the ray's color a lot. If it's clear, not so much. We're figuring out this "weight" for each point.

Here is the overview of the function:

1. Remove any opacities that are less than 0.
2. Calculate the distances between each point we queried on each ray.
3. Use the equation $\text{transparency} = 1 - \exp(-\text{density} * \text{thickness})$
4. Find the weight, which is the transparency multiplied by the cumulative product of (1-transparency) of previous segments

For step 4, we want opaque segments early in the rays to contribute to the image's final color more than segments later in the rays. (Intution: Consider looking at a wall. Do you expect to see objects behind it? Does this still hold even if the objects behind the wall are very opaque?)

```

def student_compute_weights(opacities, z_vals):
    """
    Calculate how much each point affects the ray's final color based
    on the point's density.

    Inputs:
    - opacities (torch.Tensor, [H, W, N_samples]): Information from
    our neural network about the scene's density at each point.
    - z_vals (torch.Tensor, [N_samples]): Distances along the ray
    where we checked the scene's color and density.

    Returns:
    - weights (torch.Tensor, [H,W,N_samples,4]): How much each point
    affects the ray's final color.
    """

    def raw2alpha(raw, dists, act_fn=torch.nn.functional.relu):
        return 1.0 - torch.exp(-act_fn(raw) * dists)

    # Compute 'distance' (in time) between each integration time along
    a ray.

```

```

dists = z_vals[..., 1:] - z_vals[..., :-1]

# The 'distance' from the last integration time is infinity.
# temp = torch.from_numpy(np.array([dists,
torch.broadcast_to(torch.from_numpy(np.array([1e10], dtype=float)),
dists[..., :1].shape)], dtype=np.float64))

# dists = torch.concat(
#     temp,
#     axis=-1) # [N_rays, N_samples]
dists = torch.cat( \
[dists, torch.ones_like(dists[..., 0:1]* 1e10)], \
dim=-1) # [N_rays, N_samples]

# Multiply each distance by the norm of its corresponding
direction ray
# to convert to real world distance (accounts for non-unit
directions).
# dists = dists * torch.linalg.norm(rays_d_test[..., None, :],
axis=-1) #CHANGE later

# Extract RGB of each sample position along each ray.
rgb = torch.sigmoid(opacities) # [N_rays, N_samples, 3]

# Predict density of each sample along each ray. Higher values
imply
# higher likelihood of being absorbed at this point.
alpha = raw2alpha(opacities , dists) # [N_rays, N_samples]

# Compute weight for RGB of each sample along each ray. A
cumprod() is
# used to express the idea of the ray not having reflected up to
this
# sample yet.
# [N_rays, N_samples]
T = student_cumprod(1.-alpha + 1e-10, exclusive = False)
T = torch.roll(T, 1, -1)
T[..., 0] = 1.0
weights = T*alpha
# weights = alpha *
torch.cumprod(torch.cat([torch.ones((alpha.shape[0], 1)), 1.-alpha +
1e-10], -1), -1)[:, :-1]

return weights

# delta = z_vals[..., 1:] - z_vals[..., :-1]
# delta = torch.cat([delta, 1e10 * torch.ones_like(delta[...
0:1])], dim=-1) # (height, width, samples_num) / (rays_num,

```

```

samples_num)

# # from density to transparency
# opacities = nn.functional.relu(opacities)
# alpha = 1.0 - torch.exp(-opacities * delta)

# # accumulated transmittance
# T = torch.cumprod(1 - alpha + 1e-10, dim=-1) # tricky,
(height, width, samples_num) / (rays_num, samples_num)
#s #T = student_cumprod(1 - alpha + 1e-10, exclusive=True)
# T = torch.roll(T, 1, -1)
# T[..., 0] = 1.0

# weights = T * alpha
# return weights

```

1.2.4 Putting it all Together

Using the past three functions, we will paint/render our rays to produce an image. Follow the comments in the function to create our algorithm.

```

def student_render_rays(network_fn, embedding_fn, rays_o, rays_d,
near, far, N_samples, rand=False):
    """
    Shine the rays through our scene to get a picture.

    Inputs:
    - network_fn (function): Our neural network.
    - embedding_fn (function): The positional encoding function.
    - rays_o (torch.Tensor, [H,W,3]): Where each ray of light starts.
    - rays_d (torch.Tensor, [H,W,3]): The direction each ray is
    shining.
    - near (float): How close to start looking along the ray.
    - far (float): How far to stop looking.
    - N_samples (int): How many points to check along each ray.
    - rand (bool): If true, jiggle the points a bit for smoother
    results.

    Returns:
    - rgb_map (torch.Tensor, [H,W,3]): The final color of each ray.
    - depth_map (torch.Tensor, [H,W]): The final depth of each ray.
    - acc_map (torch.Tensor, [H,W]): The final transparency of each
    ray.
    """
    # 1. Compute 3D query points along the rays.
    pts, z_vals = student_compute_query_points(rays_o, rays_d, near,
far, N_samples, rand)

    # 2. Query the neural network for each point.
    raw = student_query_network(network_fn, embedding_fn, pts)

```

```

# 3. Extract color from the network's output.
rgb = torch.sigmoid(raw[..., :3])

# 4. Compute weights for blending colors.
weights = student_compute_weights(raw[..., -1], z_vals)

# Estimated depth map is expected distance.
depth_map = torch.sum(weights * z_vals, dim=-1)

# 5. Blend the colors using the weights to get the final color for
each ray.
rgb_map = torch.sum(weights[..., None] * rgb, dim=-2) # [N_rays,
3]
acc_map = torch.sum(weights, dim=-1)

return rgb_map, depth_map, acc_map

def network_fn_test(input):
    matrix = torch.ones((3,4))
    return input @ matrix

def embedding_fn_test(input):
    return input

device = "cpu"
rays_o_test = torch.load('sol/rays_o.pt')
rays_d_test = torch.load('sol/rays_d.pt')
near_test = 0.5
far_test = 10.
N_samples_test = 64
rand_test = False
rgb_sol = torch.load('sol/rgb_map.pt')
depth_sol = torch.load('sol/depth_map.pt')
acc_sol = torch.load('sol/acc_map.pt')

check_code_correctness(student_render_rays, network_fn_test, embedding_f
n_test, \
                        rays_o_test, rays_d_test, near_test, far_test, \
N_samples_test, rand_test, test_outs=[rgb_sol, depth_sol, \
acc_sol])

True

```

2 The Neural Network: Our Magic Ray Painter

The Neural Network of NeRF is very simple. It is just a few linear layers and ReLU with skip connections. You will implement the neural network portion in this section. We will provide the training code. However, you are encouraged to modify the architecture, as we give you full points as long as you reach 22 PSNR with positional encoding.

Here is the architecture of NeRF's simple model:

- 8 intermediate layers w/ ReLU's after each
- Intermediate layers have a width of 256
- 4 output channels
- Skip connections every 4 layers (Input is fed into this layer as well as the previous layer's output)
- Input is [..., 3+3*2*L] (L is the dimensional expansion from the encoding function)

Note: To read more about building up a Neural Network, read the following article:

https://pytorch.org/tutorials/beginner/basics/buildmodel_tutorial.html

2.1.1 Create your NeRF Model

```
class student_NeRF_Model(nn.Module):
    def __init__(self, D=8, W=256, output_ch=4, skip=4, L=6):
        """
        NeRF's model.

        Input:
        D: number of layers for density
        W: number of hidden units in each layer
        skip: int that represents which layers have residuals
        concatenated to inputs
        L: size of positional encoding dimension
        """
        # super().__init__()
        # self.D = D
        # self.W = W
        # self.input_ch = 3 + 3*2*L
        # self.output_ch = output_ch
        # self.skip = skip
        # layers = []

        # self.pts_linears = nn.ModuleList(
        #     [nn.Linear(self.input_ch, W)] + [nn.Linear(W, W) if i
        not in self.skips else nn.Linear(W + self.input_ch, W) for i in
        range(D-1)])

        # ### Implementation according to the official code release
        # (https://github.com/bmild/nerf/blob/master/run_nerf_helpers.py#L104-L105)
```

```

        # # self.views_linears =
nn.ModuleList([nn.Linear(self.input_ch_views + W, W//2)])

        # ### Implementation according to the paper
        # # self.views_linears = nn.ModuleList(
        # #     [nn.Linear(input_ch_views + W, W//2)] +
[nn.Linear(W//2, W//2) for i in range(D//2)])

        # self.output_linear = nn.Linear(W, output_ch)

        # self.model = nn.ModuleList(layers)
super().__init__()
self.D = D
self.W = W
self.input_ch = 3 + 3*2*L
self.output_ch = output_ch
self.skip = skip

        layers = []
        layers.append(
            [nn.Linear(self.input_ch, self.W)] +
            [nn.Linear(self.W + self.input_ch, self.W) if i ==
self.skip \
                else nn.Linear(self.W, self.W) for i in range(self.D
- 1)]
        )

        self.activation = nn.ReLU()
        self.model = nn.ModuleList([nn.Linear(self.input_ch, self.W)]
+
            [nn.Linear(self.W + self.input_ch, self.W) if i ==
self.skip \
                else nn.Linear(self.W, self.W) for i in
range(self.D)])
        self.out = nn.Linear(self.W, self.output_ch)

        def forward(self, x):
            """
            Inputs:
            x: query inputs [B, 3+(3*2*L)] <- (The second dim is from
positional_encoding)

            Outputs:
            raws: raw outputs from model containing rgb and density for
each queried point [B, 4]
            """
            x_input = x
            for i, layer in enumerate(self.model):
                x = self.activation(layer(x))
                if i == self.skip:

```



```

        x = torch.cat([x, x_input], dim=-1)

    x = self.out(x)
    return x

```

2.1.2 Training w/o Positional Encoding

Here, you will train without positional encoding. Your model should not perform that well, as it is missing information provided by the positional encoding.

```

def no_positional_encoding(x, L=6):
    """
    Apply positional encoding to the input.

    Input:
    x: input coordinates [B,3]
    L: number of layers

    Output:
    pos_enc: Positional encoding of shape [B,3+(3*2*L)]
    """
    return x.repeat(1,13)

import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm
import math

if not os.path.exists('tiny_nerf_data.npz'):
    !wget
    http://cseweb.ucsd.edu/~viscomp/projects/LF/papers/ECCV20/nerf/tiny_nerf_data.npz

data = np.load('tiny_nerf_data.npz')
images = data['images']
poses = data['poses']
focal = data['focal']

train_images = np.concatenate([images[:101], images[102:]], axis=0)
train_poses = np.concatenate([poses[:101], poses[102:]], axis=0)
test_image = images[101]
test_pose = poses[101]

images_tensor = torch.from_numpy(train_images)
poses_tensor = torch.from_numpy(train_poses)
test_image = torch.from_numpy(test_image)
test_pose = torch.from_numpy(test_pose)

print(images_tensor.shape)
print(poses_tensor.shape)

```

```

print(test_image.shape)
print(test_pose.shape)

torch.Size([105, 100, 100, 3])
torch.Size([105, 4, 4])
torch.Size([100, 100, 3])
torch.Size([4, 4])

# Choose cuda if you have a gpu
# device = "cpu"
device = "cuda:0"

NeRF = student_NeRF_Model().to(device)
optimizer = torch.optim.Adam(NeRF.parameters(), lr=8e-4)

N_iters = 500
psnrs = []
iternums = []
i_plot = 25
# Set this lower if you get Cuda OOM
N_samples = 64

from time import time


for i in range(N_iters):
    optimizer.zero_grad()
    idx = np.random.randint(0, train_images.shape[0])
    img = images_tensor[idx]
    pose = poses_tensor[idx]
    H = img.shape[0]
    W = img.shape[1]
    rays_o, rays_d = student_get_rays(H, W, focal, pose)
    rgb, depth, acc = student_render_rays(NeRF,
no_positional_encoding, rays_o, rays_d, near=2., far=6.,
N_samples=N_samples, rand=True)
    loss = torch.mean((rgb - img)**2)
    loss.backward()
    optimizer.step()
    if i%i_plot == 0:
        H = test_image.shape[0]
        W = test_image.shape[1]
        rays_o, rays_d = student_get_rays(H, W, focal, test_pose)
        rgb, depth, acc = student_render_rays(NeRF,
no_positional_encoding, rays_o, rays_d, near=2., far=6.,
N_samples=N_samples, rand=True)
        loss = F.mse_loss(rgb, test_image)

        print(f'Iteration: {i}, Loss: {loss.item()}')
        psnr = -10. * torch.math.log(loss) / torch.math.log(10.)

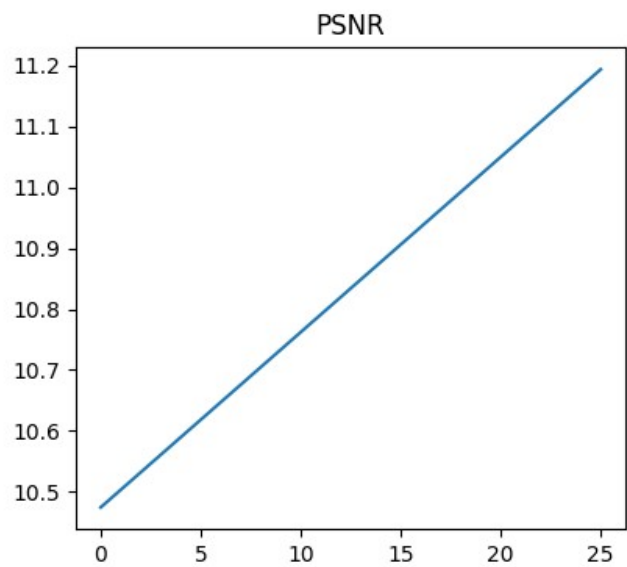
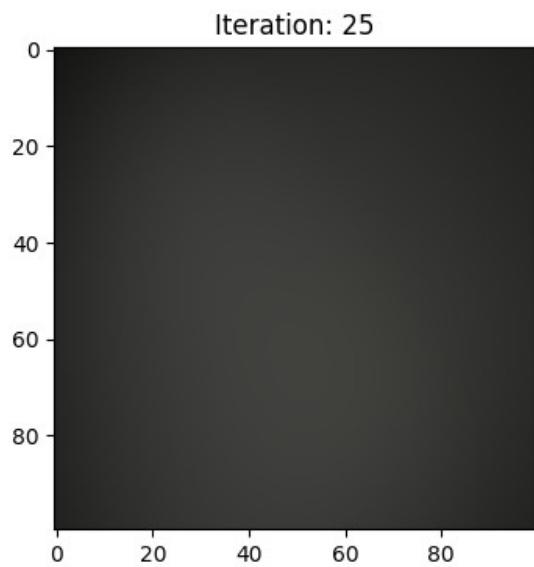
```

```
psnrs.append(psnr)
iternums.append(i)
plt.figure(figsize=(10,4))
plt.subplot(121)
plt.imshow(rgb.detach().cpu().numpy())
plt.title(f'Iteration: {i}')
plt.subplot(122)
plt.plot(iternums, psnrs)
plt.title('PSNR')
plt.show()
```



Iteration: 0, Loss: 0.08965218812227249





Iteration: 25, Loss: 0.07596400380134583




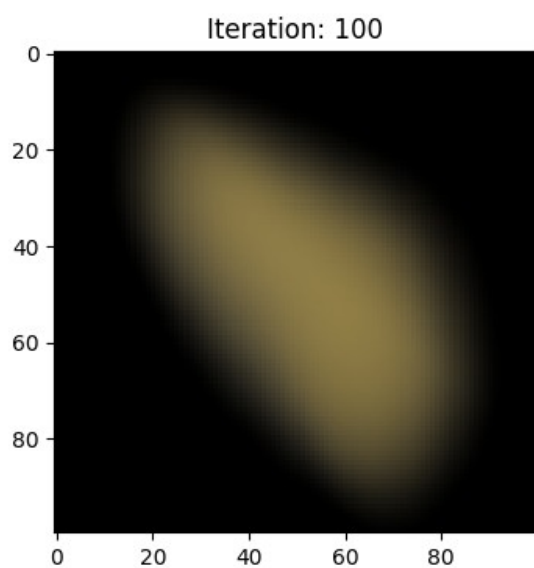
Iteration: 50, Loss: 0.059310730546712875





Iteration: 75, Loss: 0.042290251702070236





Iteration: 100, Loss: 0.04749463126063347





Iteration: 125, Loss: 0.03818003460764885





Iteration: 150, Loss: 0.043254394084215164





Iteration: 175, Loss: 0.03394508361816406





Iteration: 200, Loss: 0.06616013497114182





Iteration: 225, Loss: 0.035213686525821686





Iteration: 250, Loss: 0.031125664710998535





Iteration: 275, Loss: 0.02922290563583374





Iteration: 300, Loss: 0.02826818823814392





Iteration: 325, Loss: 0.025851847603917122





Iteration: 350, Loss: 0.027741814032197





Iteration: 375, Loss: 0.027694355696439743





Iteration: 400, Loss: 0.03539953753352165





Iteration: 425, Loss: 0.02024848759174347



Iteration: 450, Loss: 0.02225397899746895



Iteration: 475, Loss: 0.05916295573115349



2.1.3 Training w/ Positional Encoding

Here, you will train with positional encoding. You should do much better than the model without positional encoding. You are required to reach at least 22 PSNR while validating this model.

```
def positional_encoding(x, L=6):
    """
    Apply positional encoding to the input.

    Input:
    x: input coordinates [B,3]
    L: number of layers

    Output:
    pos_enc: Positional encoding of shape [B,3+(3*2*L)]
    """
    pos_enc = [x]
    for i in range(L):
        for fn in [torch.sin, torch.cos]:
            pos_enc.append(fn(2.**i * x))
    return torch.cat(pos_enc, dim=-1)

# Choose cuda if you have a gpu
# device = "cpu"
device = "cuda:0"

NeRF = student_NeRF_Model().to(device)
optimizer = torch.optim.Adam(NeRF.parameters(), lr=8e-4)

N_iters = 1000
psnrs = []
iternums = []
i_plot = 25
```

```

# Set this lower if you get Cuda OOM
N_samples = 64
threshold_reached = False

from time import time
np.random.seed(69)
torch.manual_seed(69)
for i in range(N_iters):
    optimizer.zero_grad()
    idx = np.random.randint(0, train_images.shape[0])
    img = images_tensor[idx]
    pose = poses_tensor[idx]
    H = img.shape[0]
    W = img.shape[1]
    rays_o, rays_d = student_get_rays(H, W, focal, pose)
    rgb, depth, acc = student_render_rays(NeRF, positional_encoding,
    rays_o, rays_d, near=2., far=6., N_samples=N_samples, rand=True)
    loss = torch.mean((rgb - img)**2)
    loss.backward()
    optimizer.step()
    if i%i_plot == 0:
        H = test_image.shape[0]
        W = test_image.shape[1]
        rays_o, rays_d = student_get_rays(H, W, focal, test_pose)
        rgb, depth, acc = student_render_rays(NeRF,
        positional_encoding, rays_o, rays_d, near=2., far=6.,
        N_samples=N_samples, rand=True)
        loss = F.mse_loss(rgb, test_image)



        print(f'Iteration: {i}, Loss: {loss.item()}')
        psnr = -10. * torch.math.log(loss) / torch.math.log(10.)
        if psnr > 22.0:
            threshold_reached = True

        psnrs.append(psnr)
        iternums.append(i)
        plt.figure(figsize=(10,4))
        plt.subplot(121)
        plt.imshow(rgb.detach().cpu().numpy())
        plt.title(f'Iteration: {i}')
        plt.subplot(122)
        plt.plot(iternums, psnrs)
        plt.title('PSNR')
        plt.show()



if threshold_reached:
    print('Model reached 22 PSNR!')
else:
    print('Model did not reach 22 PSNR')

```



Iteration: 0, Loss: 0.08414142578840256





Iteration: 25, Loss: 0.07058794796466827





Iteration: 50, Loss: 0.034294020384550095





Iteration: 75, Loss: 0.03376012295484543





Iteration: 100, Loss: 0.031537096947431564





Iteration: 125, Loss: 0.029409460723400116





Iteration: 150, Loss: 0.03069479949772358





Iteration: 175, Loss: 0.018764827400445938





Iteration: 200, Loss: 0.021294711157679558





Iteration: 225, Loss: 0.016696536913514137





Iteration: 250, Loss: 0.01560224499553442





Iteration: 275, Loss: 0.012983056716620922





Iteration: 300, Loss: 0.017460769042372704





Iteration: 325, Loss: 0.017315885052084923





Iteration: 350, Loss: 0.011434882879257202





Iteration: 375, Loss: 0.012903863564133644





Iteration: 400, Loss: 0.011050705797970295





Iteration: 425, Loss: 0.013505500741302967





Iteration: 450, Loss: 0.010055292397737503





Iteration: 475, Loss: 0.009111108258366585





Iteration: 500, Loss: 0.008598140440881252





Iteration: 525, Loss: 0.008741402067244053





Iteration: 550, Loss: 0.00920803938060999





Iteration: 575, Loss: 0.008204787038266659





Iteration: 600, Loss: 0.008398529142141342





Iteration: 625, Loss: 0.008130949921905994





Iteration: 650, Loss: 0.008832814171910286





Iteration: 675, Loss: 0.008265805430710316





Iteration: 700, Loss: 0.007669350132346153





Iteration: 725, Loss: 0.007064820732921362





Iteration: 750, Loss: 0.006955347489565611





Iteration: 775, Loss: 0.0074743349105119705





Iteration: 800, Loss: 0.008389628492295742





Iteration: 825, Loss: 0.007049758918583393





Iteration: 850, Loss: 0.007508496288210154





Iteration: 875, Loss: 0.007656640838831663





Iteration: 900, Loss: 0.007329396903514862





Iteration: 925, Loss: 0.008811939507722855



Iteration: 950, Loss: 0.006227308884263039



Iteration: 975, Loss: 0.006540865171700716



Model reached 22 PSNR!

2.2.1 The Fruits of our Work

Now that we have a well performing model, we can create a demo that can take in any camera position and produce the resulting scene. Great work!

```
%matplotlib inline
from ipywidgets import interactive, widgets
```

```

trans_t = lambda t : torch.tensor([
    [1,0,0,0],
    [0,1,0,0],
    [0,0,1,t],
    [0,0,0,1],
], dtype=torch.float32)

rot_phi = lambda phi : torch.tensor([
    [1,0,0,0],
    [0,np.cos(phi),-np.sin(phi),0],
    [0,np.sin(phi), np.cos(phi),0],
    [0,0,0,1],
], dtype=torch.float32)

rot_theta = lambda th : torch.tensor([
    [np.cos(th),0,-np.sin(th),0],
    [0,1,0,0],
    [np.sin(th),0, np.cos(th),0],
    [0,0,0,1],
], dtype=torch.float32)

def pose_spherical(theta, phi, radius):
    c2w = trans_t(radius)
    c2w = rot_phi(phi/180.*np.pi) @ c2w
    c2w = rot_theta(theta/180.*np.pi) @ c2w
    c2w = torch.tensor([[-1,0,0,0],[0,0,1,0],[0,1,0,0],
    [0,0,0,1]]).float() @ c2w
    return c2w

def f(**kwargs):
    c2w = pose_spherical(**kwargs)
    rays_o, rays_d = student_get_rays(H, W, focal, c2w[:3,:4])
    with torch.no_grad():
        rgb, depth, acc = student_render_rays(NeRF,
        positional_encoding, rays_o, rays_d, near=2., far=6.,
        N_samples=N_samples)
        img = np.clip(rgb,0,1)

    plt.figure(2, figsize=(20,6))
    plt.imshow(img)
    plt.show()

sldr = lambda v, mi, ma: widgets.FloatSlider(
    value=v,
    min=mi,
    max=ma,
    step=.01,

```

```
)  
names = [  
    ['theta', [100., 0., 360]],  
    ['phi', [-30., -90, 0]],  
    ['radius', [4., 2.5, 7.5]],  
]  
  
interactive_plot = interactive(f, **{s[0] : sldr(*s[1]) for s in  
names})  
output = interactive_plot.children[-1]  
output.layout.height = '350px'  
interactive_plot  
  
{"model_id": "2dbb73ddf34c4c97a4814182d0f99a54", "version_major": 2, "version_minor": 0}
```